

MACHINE LEARNING

LAB ASSESSMENT – VI

15BCE0082 VOLETI RAVI

Implement Gaussian Mixture Model Using the Expectation Maximization.

CODE:

```
import numpy as np
```

```
class GMM:
```

```
    def __init__(self, k = 3, eps = 0.0001):
```

```
        self.k = k ## number of clusters
```

```
        self.eps = eps ## threshold to stop `epsilon`
```

```
        # All parameters from fitting/learning are kept in a named tuple
```

```
        from collections import namedtuple
```

```
    def fit_EM(self, X, max_iters = 1000):
```

```
        # n = number of data-points, d = dimension of data points
```

```
        n, d = X.shape
```

```
        # randomly choose the starting centroids/means
```

```
        ## as 3 of the points from datasets
```

```
        mu = X[np.random.choice(n, self.k, False), :]
```

```

# initialize the covariance matrices for each gaussians
Sigma= [np.eye(d)] * self.k

# initialize the probabilities/weights for each gaussians
w = [1./self.k] * self.k

# responsibility matrix is initialized to all zeros
# we have responsibility for each of n points for each of k gaussians
R = np.zeros((n, self.k))

### log_likelihoods
log_likelihoods = []

P = lambda mu, s: np.linalg.det(s) ** -.5 ** (2 * np.pi) ** (-X.shape[1]/2.) \
    * np.exp(-.5 * np.einsum('ij, ij -> i', \
        X - mu, np.dot(np.linalg.inv(s), (X - mu).T).T))

# Iterate till max_iters iterations
while len(log_likelihoods) < max_iters:

    # E - Step

    ## Vectorized implementation of e-step equation to calculate the
    ## membership for each of k -gaussians
    for k in range(self.k):
        R[:, k] = w[k] * P(mu[k], Sigma[k])

    ### Likelihood computation

```

```

log_likelihood = np.sum(np.log(np.sum(R, axis = 1)))

log_likelihoods.append(log_likelihood)

## Normalize so that the responsibility matrix is row stochastic
R = (R.T / np.sum(R, axis = 1)).T

## The number of datapoints belonging to each gaussian
N_ks = np.sum(R, axis = 0)

# M Step

## calculate the new mean and covariance for each gaussian by
## utilizing the new responsibilities
for k in range(self.k):

    ## means
    mu[k] = 1. / N_ks[k] * np.sum(R[:, k] * X.T, axis = 1).T
    x_mu = np.matrix(X - mu[k])

    ## covariances
    Sigma[k] = np.array(1 / N_ks[k] * np.dot(np.multiply(x_mu.T, R[:, k]), x_mu))

    ## and finally the probabilities
    w[k] = 1. / n * N_ks[k]

# check for onvergence
if len(log_likelihoods) < 2 : continue
if np.abs(log_likelihood - log_likelihoods[-2]) < self.eps: break

```

```

## bind all results together

from collections import namedtuple

self.params = namedtuple('params', ['mu', 'Sigma', 'w', 'log_likelihoods', 'num_iters'])

self.params.mu = mu

self.params.Sigma = Sigma

self.params.w = w

self.params.log_likelihoods = log_likelihoods

self.params.num_iters = len(log_likelihoods)


return self.params


def plot_log_likelihood(self):

    import pylab as plt

    plt.plot(self.params.log_likelihoods)

    plt.title('Log Likelihood vs iteration plot')

    plt.xlabel('Iterations')

    plt.ylabel('log likelihood')

    plt.show()


def predict(self, x):

    p = lambda mu, s : np.linalg.det(s) ** -0.5 * (2 * np.pi) ** \
        (-len(x)/2) * np.exp( -0.5 * np.dot(x - mu , \
            np.dot(np.linalg.inv(s) , x - mu)))

    probs = np.array([w * p(mu, s) for mu, s, w in \
        zip(self.params.mu, self.params.Sigma, self.params.w)])

    return probs/np.sum(probs)

```

```

def demo_2d():

    # Load data

    #X = np.genfromtxt('data1.csv', delimiter=',')

    ### generate the random data

    np.random.seed(3)

    m1, cov1 = [9, 8], [[.5, 1], [.25, 1]] ## first gaussian
    data1 = np.random.multivariate_normal(m1, cov1, 90)

    m2, cov2 = [6, 13], [[.5, -.5], [-.5, .1]] ## second gaussian
    data2 = np.random.multivariate_normal(m2, cov2, 45)

    m3, cov3 = [4, 7], [[0.25, 0.5], [-0.1, 0.5]] ## third gaussian
    data3 = np.random.multivariate_normal(m3, cov3, 65)
    X = np.vstack((data1, np.vstack((data2, data3))))
    np.random.shuffle(X)

#   np.savetxt('sample.csv', X, fmt = "%.4f", delimiter = ",")

    #####

    gmm = GMM(3, 0.000001)
    params = gmm.fit_EM(X, max_iters= 100)
    print params.log_likelihoods

    import pylab as plt

    from matplotlib.patches import Ellipse

    def plot_ellipse(pos, cov, nstd=2, ax=None, **kwargs):

        def eigsorted(cov):

            vals, vecs = np.linalg.eigh(cov)

```

```

        order = vals.argsort()[::-1]

        return vals[order], vecs[:,order]

if ax is None:
    ax = plt.gca()

vals, vecs = eigsorted(cov)

theta = np.degrees(np.arctan2(*vecs[:,0][::-1]))

# Width and height are "full" widths, not radius
width, height = 2 * nstd * np.sqrt(abs(vals))

ellip = Ellipse(xy=pos, width=width, height=height, angle=theta, **kwargs)

ax.add_artist(ellip)

return ellip

def show(X, mu, cov):

    plt.cla()

    K = len(mu) # number of clusters

    colors = ['b', 'k', 'g', 'c', 'm', 'y', 'r']

    plt.plot(X.T[0], X.T[1], 'm*')

    for k in range(K):

        plot_ellipse(mu[k], cov[k], alpha=0.6, color = colors[k % len(colors)])

fig = plt.figure(figsize = (13, 6))

fig.add_subplot(121)

```

```
show(X, params.mu, params.Sigma)
fig.add_subplot(122)
plt.plot(np.array(params.log_likelihoods))
plt.title('Log Likelihood vs iteration plot')
plt.xlabel('Iterations')
plt.ylabel('log likelihood')
plt.show()
print gmm.predict(np.array([1, 2]))
```

```
if __name__ == "__main__":
```

```
    demo_2d()
```

```
    from optparse import OptionParser
```

```
    parser = OptionParser()
```

```
    parser.add_option("-f", "--file", dest="filepath", help="File path for data")
```

```
    parser.add_option("-k", "--clusters", dest="clusters", help="No. of gaussians")
```

```
    parser.add_option("-e", "--eps", dest="epsilon", help="Epsilon to stop")
```

```
    parser.add_option("-m", "--maxiters", dest="max_iters", help="Maximum no. of iteration")
```

```
    options, args = parser.parse_args()
```

```
    if not options.filepath : raise('File not provided')
```

```
    if not options.clusters :
```

```
        print("Used default number of clusters = 3" )
```

```
        k = 3
```

```
    else: k = int(options.clusters)
```

```
if not options.epsilon :  
    print("Used default eps = 0.0001" )  
    eps = 0.0001  
else: eps = float(options.epsilon)  
  
if not options.max_iters :  
    print("Used default maxiters = 1000" )  
    max_iters = 1000  
else: eps = int(options.maxiters)  
  
X = np.genfromtxt(options.filepath, delimiter=',')  
gmm = GMM(k, eps)  
params = gmm.fit_EM(X, max_iters)  
print params.log_likelihoods  
gmm.plot_log_likelihood()  
print gmm.predict(np.array([1, 2]))
```


OUTPUT:

```
xelese@xelese-Lenovo-Y50-70: ~/Machine Learning/GMM
xelese@xelese-Lenovo-Y50-70:~/Machine Learning/GMM$ python GMM.py
GMM.py:113: RuntimeWarning: covariance is not positive-semidefinite.
  data1 = np.random.multivariate_normal(m1, cov1, 90)
GMM.py:116: RuntimeWarning: covariance is not positive-semidefinite.
  data2 = np.random.multivariate_normal(m2, cov2, 45)
GMM.py:119: RuntimeWarning: covariance is not positive-semidefinite.
  data3 = np.random.multivariate_normal(m3, cov3, 65)
[-931.457086088967, -496.48852371740338, -459.33914688320544, -452.0181991849052
6, -451.17650595193993, -448.73153830461285, -428.21188748537418, -373.731963991
36657, -296.38814029334884, -184.92988425156602, -145.12717607796412, -143.94873
831421006, -143.94886545905223, -143.9488655473433]
```

